

# Assignment 4: Optimizing Matrix Operations for Better Cache Locality (100 points)

This assignment has two parts.

## Shared Structure of Both Parts

Both parts re-use much of the same code, so the general instructions for how the code is laid out and how to edit/grade your submission is listed here.

### General Information

- All matrices in the code are implemented as **1D arrays** that are indexed by multiplying the row offset by the number of values in a column. Can you think of why this is better for cache locality?
- The function `int *access_matrix(matrix_t *mat, int row, int col)` returns a pointer to the integer in a `matrix_t` at row `row` and column `col`. You can read/write this value by dereferencing.

### Editing

**Only edit the code in the part of the file labelled:**

```
// ONLY EDIT THIS LOOP
```

- Importantly, do **not** touch the `MARKER_START` and `MARKER_END` variables in the code. **You may only edit code between the two `MARKER` variables, as this is where the implemented algorithms are.**
- Do **not** touch anything in the `res` folder. These files are required for the autograder to function.

### Grading

- There are 5 test cases. A program that yields a valid matrix gives you 30% of the points for the test case. A program that has less cache misses and evictions than the original gives you 70% of the points for the test case.
- To test your program, simply run `./autograder`.

### Testing

You cannot test the cache yourself, as this involves filtering the file generated by our cache simulator, `csim-ref`, but you may compile your program yourself to make sure your matrix operations are working properly. To do this:

```
gcc -Wall -O0 -g matadd.c/matmul.c -o matadd/matmul
```

Choose the appropriate (`matadd/matmul`) depending on the part you are working on.

**Running `matadd`** Choose your matrix dimension (suppose it is `n x n`) and choose your values for the matrix. Suppose we choose a `2 x 2` matrix with values

```
1 2
31 3
```

Then we would invoke `./matadd n n entries of n`, which is `./matadd 2 2 1 2 31 3` in this case.

**Running `matmul`** Choose your matrix dimension (suppose it is  $n \times n$ ) and choose your values for the matrix. Suppose we choose a  $2 \times 2$  matrix with values

```
1 2
31 3
```

And a right matrix with values

```
38 1
1 2
```

Then we would invoke `./matmul n n n entries of left entries of right`, which is `./matmul 2 2 2 1 2 31 3 38 1 1 2` in this case. Note how we just put the right matrix's values immediately after the left matrix.

**In both cases, the output result is stored in the `.mat` file.** You can view this file however you want. In the terminal, for example, you may invoke `cat .mat` after running `matmul` or `matadd` to see the resulting matrix.

## Assumptions

- All input matrices are square.
- The cache in both parts is a 1-way (direct mapped) cache with 4 sets, and the number of addresses in each block is 4. An LRU eviction policy is used.

## Part 1: Matrix Addition (`matadd`) (40 points)

In this part, you will edit `matadd.c` to optimize adding 1 to every entry in a matrix. The current code written is inefficient from a spatial locality standpoint.

Currently, the code adds 1 to every entry matrix by looping through every **row** entry per column. Why does this program yield bad spatial locality as a result?

**Optimize the code by editing the indicated sections as discussed in the previous section, and then run the autograder to make sure your code works.**

## Part 2: Matrix Multiplication (`matmul`) (60 points)

In this part, you will edit `matmul.c` to improve cache locality in matrix multiplication. Note that the normal pattern of accessing the matrix is not the most cache-optimal way to multiply a matrix, as the right matrix in the matrix multiplication is accessed in a way that is not conducive to spatial locality. Can you change the loop in the `matmul` function to be more optimal to a cache?

**Optimize the code by editing the indicated sections as discussed in the previous section, and then run the autograder to make sure your code works.**

## Hints

You may consult this slide deck (<https://rutgers.instructure.com/files/40122456/>) to help with your understanding.

The file is also on Canvas: Files -> A4 -> `a4_slides.pdf`.

## Submission

To submit your code, simply invoke `tar cvf a4.tar a4/` and submit the `a4.tar` file.

To verify your submission, you can check the contents of the tar file without extracting it using `tar tvf a4.tar`. You should see something similar (but it's okay if it's not exactly the same as)

```

drwxr-xr-x NETID/NETID    0 2024-04-22 23:34 a4/
drwxr-xr-x NETID/NETID    0 2024-04-22 23:44 a4/matmul/
-rwxr-xr-x NETID/NETID 4871768 2024-04-22 23:44 a4/matmul/autograder
-rw-r--r-- NETID/NETID    699 2024-04-22 23:44 a4/matmul/.mat
-rw-r--r-- NETID/NETID    18 2024-04-22 23:44 a4/matmul/.csim_results
-rw-r--r-- NETID/NETID    13 2024-04-22 23:44 a4/matmul/.marker
-rw-r--r-- NETID/NETID 887730 2024-04-22 23:44 a4/matmul/.trace
drwxr-xr-x NETID/NETID    0 2024-04-22 23:44 a4/matmul/res/
-rw-r--r-- NETID/NETID   2361 2024-04-22 23:39 a4/matmul/res/matmul.c
-rwxr-xr-x NETID/NETID 25592 2024-04-22 23:37 a4/matmul/res/csim-ref
-rw-r--r-- NETID/NETID   2361 2024-04-22 23:40 a4/matmul/matmul.c
drwxr-xr-x NETID/NETID    0 2024-04-22 23:45 a4/matadd/
-rw-r--r-- NETID/NETID   1828 2024-04-22 23:40 a4/matadd/matadd.c
-rwxr-xr-x NETID/NETID 4866848 2024-04-22 23:36 a4/matadd/autograder
-rw-r--r-- NETID/NETID    375 2024-04-22 23:42 a4/matadd/.mat
-rw-r--r-- NETID/NETID    14 2024-04-22 23:42 a4/matadd/.csim_results
-rw-r--r-- NETID/NETID    13 2024-04-22 23:42 a4/matadd/.marker
-rw-r--r-- NETID/NETID 35530 2024-04-22 23:42 a4/matadd/.trace
drwxr-xr-x NETID/NETID    0 2024-04-22 23:45 a4/matadd/res/
-rw-r--r-- NETID/NETID   1828 2024-04-22 23:40 a4/matadd/res/matadd.c
-rwxr-xr-x NETID/NETID 25592 2024-04-22 23:37 a4/matadd/res/csim-ref

```